

# **Reducing Static Power Consumption on Multiprocessor Real-Time systems**

Séminaire MeFoSyLoMa

Vincent Legout, Mathieu Jan, Laurent Pautet

CEA LIST, Télécom ParisTech

Mars 1, 2013

# Introduction

- Context: embedded hard real-time systems (e.g. automotive, . . . )
  - Hard real-time: time constraints cannot be violated
- Reducing energy consumption has several advantages:
  - Increase the autonomy of battery-powered systems
  - Provide energy efficient or “green” solutions
- Target multiprocessor systems
  - Uniprocessor are now superseded by multiprocessor systems
  - No power-aware optimal multiprocessor scheduling algorithm

## Objective

Schedule multiprocessor real-time systems to reduce consumption

# Modeling the consumption of processors



## Focusing on static consumption

- Dynamic and static consumption ( $P = P_{dynamic} + P_{static}$ )
  - $P_{static} = \text{constant}$  (*leakage current*)
- Static consumption now dominates dynamic consumption
  - Technology evolves: higher density, smaller supply voltage
- 2 software solutions to reduce consumption
  - Dynamic consumption: Dynamic Voltage & Frequency Scaling
  - Static consumption: Dynamic Power Management

# DPM: Dynamic Power Management

Reduce static consumption by using the low-power states of processors

- In a low-power state:
  - Processor inactive and energy consumption reduced
  - Transition delay  $\delta$  required to get back to the active state
- Several low-power states available
  - The lower the consumption, the higher the transition delay
- Break-Even Time (BET): smallest idle period for which activating a low-power state saves more energy than letting the processor idle




Objective: activate the most efficient low-power states without increasing the number of preemptions and migrations

# DPM multiprocessor scheduling algorithms

The objective being to create large idle periods

- Partitioned scheduling (no migration)
  - Use uniprocessor DPM solutions on each processor
  - Cannot merge idle periods from different processors
- Global scheduling (migrations allowed)
  - Minimize the number of processors [Bhatti09]
    - Non-optimal scheduling algorithm
- Existing DPM scheduling algorithms are not suitable
  - Schedule decisions should be taken according to the workload
  - Instead, they use task characteristics (e.g. earliest deadline)

	$\tau_1$	$\tau_2$
WCET	1	2
Period	2	3



# Approach

- Solve the problem off-line. Several advantages:
  - Compute a schedule in a hyperperiod
    - Without taking scheduling decisions for each task individually
    - But for the task set as a whole in a hyperperiod
  - Use linear programming to express constraints and objectives:
    - Real-time constraint: no deadline miss
    - Consumption objective: large idle periods
  - Ensure a maximal consumption in the worst case scenario
- Computation using the Worst Case Execution Time (WCET)
- Larger idle periods means minimizing the number of idle periods
- LPDPM: Linear Programming DPM


# Task model

- $m$  processors,  $n$  tasks (WCET  $C$ , period  $T$ ), hyperperiod  $H$
- Global utilization  $U = \sum_{i=1}^n \frac{C_i}{T_i}$  ( $m - 1 < U < m$ )

- Division in intervals, for example :

	$\tau_1$	$\tau_2$	$\tau_3$
WCET	0.8	2.4	4
Period	4	4	6
Jobs	3	3	2

- Objective to compute all weights of all tasks on all intervals
  - $w_{j,k}$  weight of job  $j$  on interval  $k$



# Initial linear system

- Set of constraints to keep the schedulability [Lemerre08]:

- $|I_k|$ : length of interval  $k$
- $J_k$ : set of jobs in interval  $k$
- $E_j$ : set of intervals where  $j$  is present

$$\forall k, \sum_{j \in J_k} w_{j,k} \leq m$$


$U \leq m$  inside each interval

$$\forall k, \forall j, 0 \leq w_{j,k} \leq 1$$

$0 \leq u \leq 1$  for each task inside each interval

$$\forall j, \sum_{k \in E_j} w_{j,k} \times |I_k| = j.c$$

All jobs are fully executed






# Add an additional idle task $\tau'$

Accounts for the time where processors are supposed to be idle

- Utilization of  $\tau' = m - U$ , *period* =  $H$
- As a consequence: global utilization  $U = m$
- Only one processor can be idle simultaneously
- $w_k$  weight of  $\tau'$  on interval  $k$
- Objective: reduce the number of preemptions of  $\tau'$

Interval	1	2	3	4
$w_k$	0.2	0.6	0.8	0.7




# Add an additional idle task $\tau'$

Accounts for the time where processors are supposed to be idle

- Utilization of  $\tau' = m - U$ , *period* =  $H$
- As a consequence: global utilization  $U = m$
- Only one processor can be idle simultaneously
- $w_k$  weight of  $\tau'$  on interval  $k$
- Objective: reduce the number of preemptions of  $\tau'$

Interval	1	2	3	4
$w_k$	0.2	0.6	0.8	0.7



# Reducing the number of preemptions of $\tau'$ (1)


Avoid intervals in which  $\tau'$  is preempted (i.e.  $0 < w_k < 1$ )

- Maximize the number of intervals for which  $w_k = 0$  and the number of intervals for which  $w_k = 1$
- 2 binary variables  $f_k$  (*full*) and  $e_k$  (*empty*) such as:

$$f_k = \begin{cases} 0 & \text{if } w_k = 1 \\ 1 & \text{otherwise} \end{cases}$$

$$e_k = \begin{cases} 0 & \text{if } w_k = 0 \\ 1 & \text{otherwise} \end{cases}$$

$$\text{Minimize } \sum_k f_k + e_k$$



# Reducing the number of preemptions of $\tau'$ (1)


Avoid intervals in which  $\tau'$  is preempted (i.e.  $0 < w_k < 1$ )

- Maximize the number of intervals for which  $w_k = 0$  and the number of intervals for which  $w_k = 1$
- 2 binary variables  $f_k$  (*full*) and  $e_k$  (*empty*) such as:

$$f_k = \begin{cases} 0 & \text{if } w_k = 1 \\ 1 & \text{otherwise} \end{cases}$$

$$e_k = \begin{cases} 0 & \text{if } w_k = 0 \\ 1 & \text{otherwise} \end{cases}$$

$$\text{Minimize } \sum_k f_k + e_k$$




# Reducing the number of preemptions of $\tau'$ (2)

Avoid preemptions of  $\tau'$  on interval boundaries

- Group full intervals and group empty intervals
- 2 binary variables  $fc_k$  and  $ec_k$  such as:

$$fc_k = \begin{cases} 1 & \text{if } f_k = 1 \text{ and } f_{k+1} = 0 \\ 0 & \text{otherwise} \end{cases} \quad ec_k = \begin{cases} 1 & \text{if } e_k = 1 \text{ and } e_{k+1} = 0 \\ 0 & \text{otherwise} \end{cases}$$

$$\text{Minimize } \sum_k f_k + e_k + fc_k + ec_k$$




# Reducing the number of preemptions of $\tau'$ (2)

Avoid preemptions of  $\tau'$  on interval boundaries

- Group full intervals and group empty intervals
- 2 binary variables  $fc_k$  and  $ec_k$  such as:

$$fc_k = \begin{cases} 1 & \text{if } f_k = 1 \text{ and } f_{k+1} = 0 \\ 0 & \text{otherwise} \end{cases} \quad ec_k = \begin{cases} 1 & \text{if } e_k = 1 \text{ and } e_{k+1} = 0 \\ 0 & \text{otherwise} \end{cases}$$


$$\text{Minimize } \sum_k f_k + e_k + fc_k + ec_k$$



# On-line schedule inside intervals

Using the weights computed off-line using the linear program


- Schedule impacts idle periods only when  $0 < w_k < 1$ 
  - Solution: schedule  $\tau'$  at the beginning or at the end the interval
- During the execution of  $\tau'$ , activate the deepest low-power state
  - Processor stays idle if the idle period is not large enough



# On-line schedule inside intervals

Using the weights computed off-line using the linear program

- Schedule impacts idle periods only when  $0 < w_k < 1$ 
  - Solution: schedule  $\tau'$  at the beginning or at the end the interval
- During the execution of  $\tau'$ , activate the deepest low-power state
  - Processor stays idle if the idle period is not large enough






# Experimental evaluation

- In a simulator, schedule random task sets with LPDPM and two existing optimal multiprocessor schedulers RUN & U-EDF
- No other DPM optimal multiprocessor scheduling algorithm
- RUN & U-EDF: schedulers reducing the number of preemptions
- 4 processors, 10 tasks
- 2000 task sets for each global utilization between 3.05 and 3.95
- Objective: evaluate the number and the length of the idle periods as well as the number of preemptions

# Mean number of idle periods


- Between 2 and 5 times less idle periods for LPDPM



# Idle period lengths


Using all idle periods generated by all task sets

- RUN & U-EDF: smaller idle periods, with a length always  $< 120$
- LPDPM: larger idle periods, can be twice as large



# Mean total number of preemptions

- Fewer preemptions than U-EDF
- Less than 1.4 more preemptions than RUN



# Conclusion and perspectives

- Contribution: LPDPM, an off-line optimal multiprocessor real-time scheduling algorithm reducing static consumption
  - LPDPM generates less and larger idle periods such that deeper low-power states can be activated. The total number of preemptions does not increase.
  - Simulation of the global consumption using the low-power states of the STM32L board (Cortex-M3): LPDPM is up to 8% more energy efficient than existing optimal scheduling algorithms
- Perspectives
  - When tasks do not use their WCET, schedule intervals to further extend the existing idle periods
  - Temperature: impact on the energy consumption